

ECS 158 Final Report:
combn() in Snow, OpenMP, and CUDA

Stefan Peterson
Bijan Agahi
Arjun Bharadwaj

Abstract

Parallel architectures are becoming increasingly prevalent in the modern world and as such have spawned multiple frameworks to assist with parallel programming, but these frameworks vary greatly in their implementation and performance. This paper takes a close look at three frameworks: R-Snow, OpenMP and CUDA. We took on the task of re-writing the `combn()` function found in R using each of these frameworks. We measure run times of these tests at different sizes and with varying inputs and levels of difficulty. In the end, we point out the strengths and weaknesses of each language as well as note on features which we, as programmers, find particularly useful or troublesome.

Contents

1 Introduction

With the relatively recent introduction of widely available multi-processor CPUs in household computers, parallel programming has become a hot topic in the world of computer science. In particular, the graphics and games industry has driven the multithreaded programming style to the forefront of everyone's attention. Through the process of splitting up tasks and distributing those tasks to individual cores on the CPU or GPU, parallel processing in most cases can provide vast improvements in both performance and efficiency. OpenMP, R-Snow, and CUDA are three such parallel languages, and this paper aims to compare and contrast the strengths, weaknesses, and performance of each language.

1.1 Setup for Benchmarks

All tests for the purposes of this paper were run on Tetraat University of California, Davis. This computer consists of a 64-bit 8-core processor with 16 GB of RAM and an Intel Core i7-2600K CPU with a clock speed of 3.4 GHz and a cache size of 8192 KB. The machine was running release 19 of the linux distribution Fedora.

2 `combn()` in R-Snow

2.1 `combn()` in R-Snow: Strategy

2.2 `combn()` in R-Snow: Implementation

2.3 `combn()` in R-Snow: Data

2.4 `combn()` in R-Snow: Analysis

3 `combn()` in OpenMP

3.1 Strategy

For the OpenMP version of the combinations function, we decided to parallelize by dedicating each thread to a subsection of the problem


3.2 Implementation

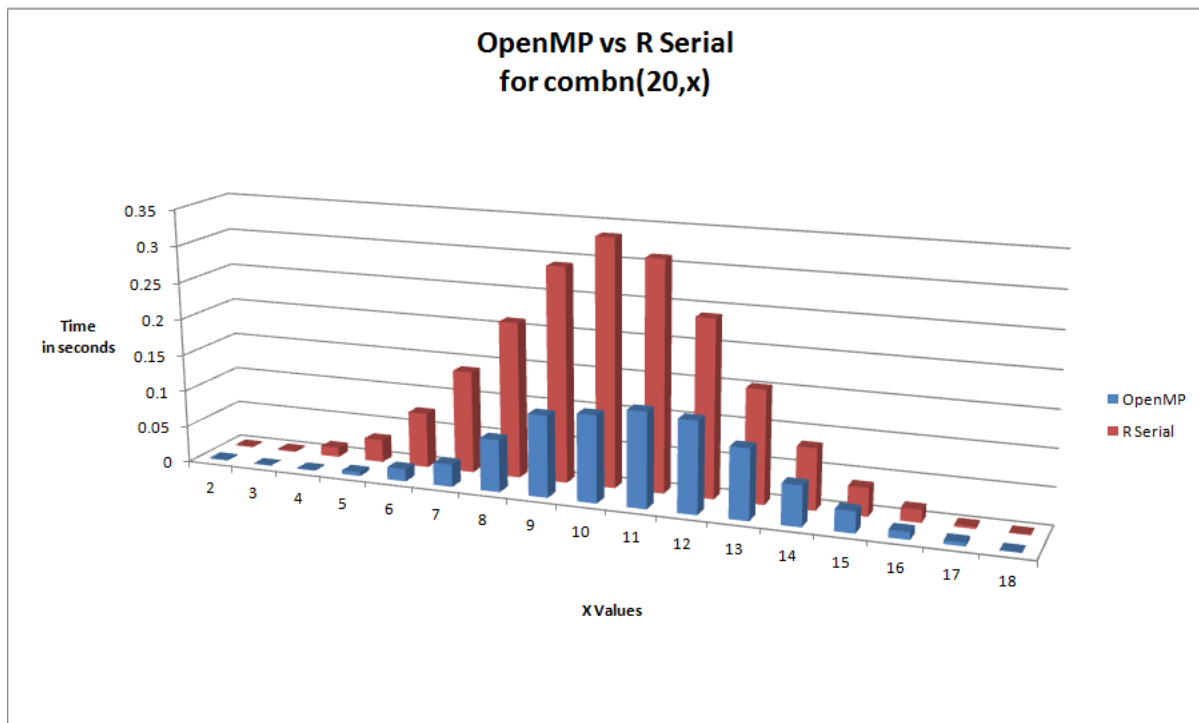
All the references to code will be references to Appendix: A.1

We decided to take a recursive approach to the implementation of the combinations function in C++. The core functionality of the program can be found on lines 23-33 in the method `findCombs`. This method takes in as its parameters 2 integers, a vector, and a 2d vector. Upon the first call of `findCombs()` the vector combination will already have been initialized to contain the first value in the combination (i.e. if we assume we are looking for the combination [1,2,3] the vector will

be initialized to contain [1] and this is done on line 71 in the `combn()` function. From there the recursive algorithm will add the next number (i.e. combinations will now contain [1,2]) and then will call `findCombs()` recursively until the base case is met upon which the vector combination is pushed on to the 2D vector. When the recursion is done and the 2D vector contains all its possible combinations, the function `vectorToArray()` is called and the 2D vector is parsed and its values are added to the appropriate region of the global array.

3.3 Data





3.4 Analysis

4 combn() in Thrust

4.1 combn() in Thrust: Strategy

4.2 combn() in Thrust: Implementation

4.3 combn() in Thrust: Data

4.4 combn() in Thrust: Analysis

5 Research Papers

5.1 Authors — Position

5.1.1 Description

5.1.2 Analysis

6 Conclusion

stuff here

7 Bibliography

A Appendix: Code

A.1 combn() — OpenMP implementation

combnOpenMP.cpp

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <iostream>
4 #include <vector>
5 #include <algorithm>    // std::find
6 #include <omp.h>
7
8 using namespace std;
9
10 vector<int> values;
11 int *allCombs;
12 int *numEntriesPerLevel;
13
14 double choose(int n, int k) {
15     if(k == 0) return 1;
16     return (n * choose(n - 1, k - 1)) / k;
17 }
18
19 void addComb(vector<vector<int> > &v, vector<int> &v2) {
20     v.push_back(v2);
21 }
22
23 void findCombs(int offset, int k, vector<vector<int> > &v, vector<int>
    combination) {
24     if (k == 0) {
25         addComb(v, combination);
26         return;
27     }
28     for(int i = offset; i <= values.size() - k; ++i) {
29         combination.push_back(values[i]);
30         findCombs(i + 1, k - 1, v, combination);
31         combination.pop_back();
32     }
33 }
34
35 void findEntriesPerLevel(int *array, int x, int m) {
36     array[0] = 0;
37     for(int i = 1; i < x - m + 1; i++) {
38         array[i] = array[i - 1] + (choose(x - i, m - 1) * m);
39     }
40 }
41
42 void vectorToArray(vector<vector<int> > &v, int a[], int startIndex) {
43     vector<vector<int> >::const_iterator row;
44     vector<int>::const_iterator col;
45     int index = startIndex;
46
47     for (row = v.begin(); row != v.end(); ++row)
48     {
49         for (col = row->begin(); col != row->end(); ++col)
50         {
51             a[index] = *col;
52             index++;
```



```

53     }
54 }
55 }
56
57 int * combn(int x, int m) {
58     int size = choose(x,m);
59     for(int i = 0; i < x; ++i)
60         values.push_back(i+1);
61
62     allCombs = new int[size * m];
63     numEntriesPerLevel = new int[x - m + 1];
64     findEntriesPerLevel(numEntriesPerLevel, x, m);
65
66
67     #pragma omp parallel for schedule(dynamic)
68     for(int i = 1; i <= x - m + 1; i++){
69         vector<vector<int> > levelCombinations;
70         vector<int> combination;
71         combination.push_back(i);
72         findCombs(i, m - 1, levelCombinations, combination);
73         vectorToArray(levelCombinations, allCombs, numEntriesPerLevel[i - 1]);
74         combination.pop_back();
75     }
76     #pragma omp barrier
77
78     return allCombs;
79 }

```

This is a recursive implementation of the combinations function. Here we used `#pragma omp parallel for schedule(dynamic)` on line 67 to help load balance the parallelization of the threading.

A.2 combn() — Thrust implementation

combnThrust.cilk

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <iostream>
4
5 using namespace std;
6
7
8
9
10 int * combn(int x, int m) {
11     int *ret = new int[x * m];
12
13     int loc = 0;
14     for(int i = 0; i < x; i++) {
15         for(int j = i; j < i + m; j++) {
16             ret[loc] = j;
17             loc++;
18         }
19     }
20     return ret;
21 }
22
23
24 template <typename T>

```

```

25 T * combn(T x, int m) {
26
27
28 }
29
30
31 int main (int argc, char** argv) {
32
33     int x = 10;
34     int m = 2;
35     int * vals[x * m];
36     vals = combn(x,m);
37     for (int i = 0; i < x * m; i++) {
38         cout << vals[x];
39     }
40     return 0;
41 }

```

This code is directly translate from the OpenMP version above, using the same algorithm. We have various Cilk++ equivalents of OpenMP, including: the

There are also a variety of files specific for testing purposes included in the submitted .tar file.

B Appendix: Member Contributions

Stefan Peterson

- L^AT_EX Formatting, editing
- OpenMP implementation
- Abstract
- Introduction
- Research paper summarization/analysis
- Conclusion
- Data analysis writeup
- Relevant research

Bijan Agahi

- R-Snow implementation
- Experiment data collection
- Test scripting
- Graphing in R, tabling
- Relevant research

Arjun Bharadwaj

- Thrust implementation
- Relevant research